

# DSA15207 Modern Deep Learning

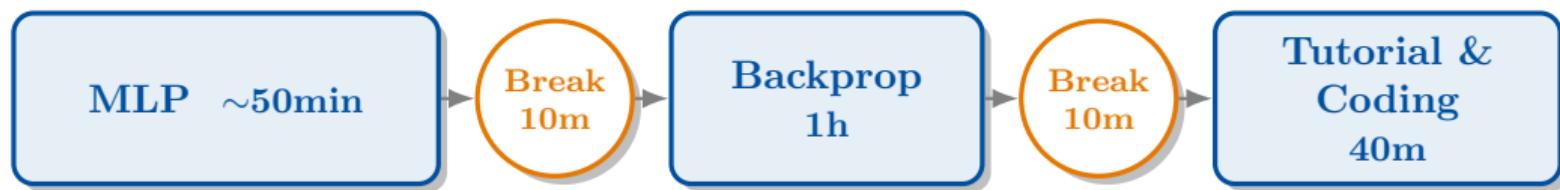
Multilayer Perceptrons (MLP)

Xingyi Yang

Department of Data Science and Artificial Intelligence

Hong Kong Polytechnic University

# Agenda for Today



# Contents



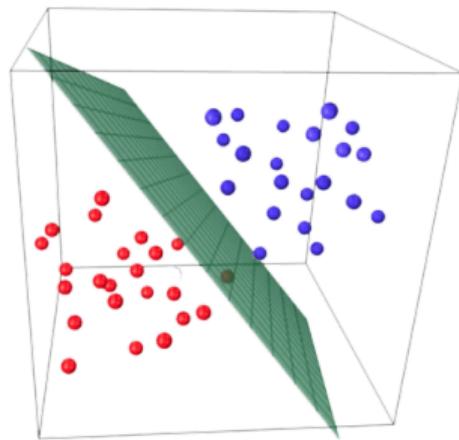
- ▶ Limitations of Linear Models
- ▶ Multilayer Perceptrons (MLPs)
- ▶ MLPs in the Modern Era

# Recap: Linear Classification



- Recall the simple linear classifier (Perceptron):

$$y = f(\mathbf{w}^T \mathbf{x} + b)$$



- It defines a linear decision boundary (hyperplane) in the input space.
- **Question:** Is a single linear unit enough for all problems?

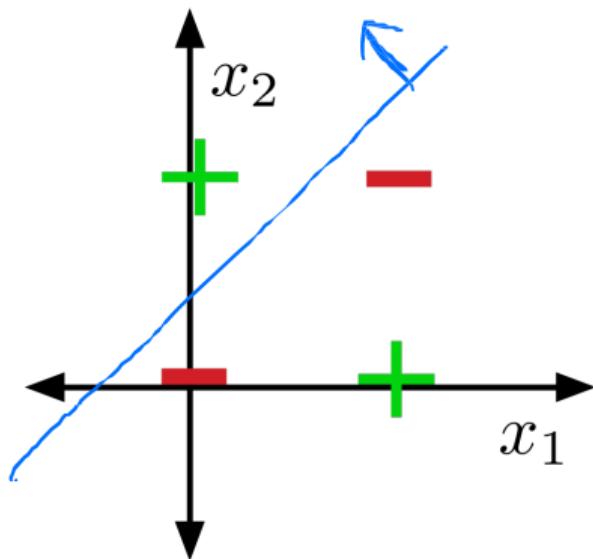
# Example I: The XOR Problem



Consider the **XOR logic function**:

| $x_1$ | $x_2$ | $y$ |
|-------|-------|-----|
| 0     | 0     | 0   |
| 0     | 1     | 1   |
| 1     | 0     | 1   |
| 1     | 1     | 0   |

Can we separate the classes 0 and 1 with a single straight line?



**No!** XOR is not linearly separable.

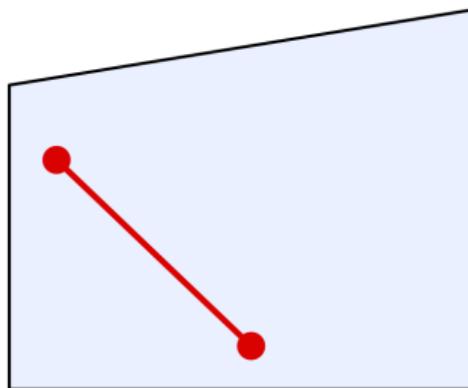
# The XOR Problem are not Linear Separable



## Convex Sets

A set  $S$  is convex if any line segment connecting points in  $S$  lies entirely within  $S$ .

$$\mathbf{x}_1, \mathbf{x}_2 \in S \implies \lambda \mathbf{x}_1 + (1 - \lambda) \mathbf{x}_2 \in S, \quad \forall \lambda \in [0, 1]$$

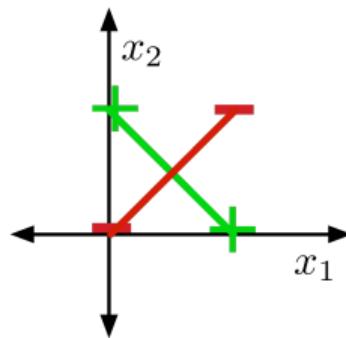


# The XOR Problem are not Linear Separable



Let's prove that XOR is not linearly separable

- Linear classifiers divide space into two **half-spaces**, which are convex sets.



- If positive examples  $(0, 1)$  and  $(1, 0)$  are in the positive half-space, the line connecting them must be too.
- If negative examples  $(0, 0)$  and  $(1, 1)$  are in the negative half-space, their connecting line must be too.
- These two lines **intersect**. intersection can't lie in both half-spaces. **Contradiction!!!!**

## Example II: Translation Invariance



Binary 16D vectors (white = 0, black = 1).

**Goal:** Distinguish patterns A and B under any cyclic translation.



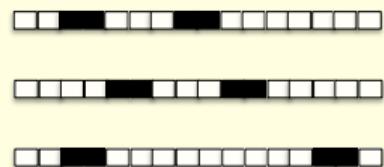
pattern A



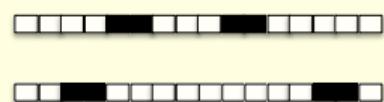
pattern A



pattern A



pattern B



pattern B



pattern B

**Key Insight:**

Linear classifier  $y = \sum w_i x_i$  uses fixed per-pixel weights

⇒ Shifting the image  $x$  activates different weights  $w$ . → output changes

⇒ To be invariant ( $y$  stays same), all weights  $w_i$  must be equal

⇒ If all weights are equal, the model **only counts the number of black pixels**

⇒ Patterns A and B (same 1 count) are indistinguishable.

**Conclusion:** Linear classifiers **cannot be translation invariant!**

# Overcoming Limits: Feature Maps



- We can solve **XOR** by mapping input to a higher-dimensional space  $\psi(\mathbf{x})$ .

# Overcoming Limits: Feature Maps



- We can solve **XOR** by mapping input to a higher-dimensional space  $\psi(\mathbf{x})$ .
- Example:  $\psi(\mathbf{x}) = [x_1, x_2, x_1x_2]^T$ .

# Overcoming Limits: Feature Maps



- We can solve **XOR** by mapping input to a higher-dimensional space  $\psi(\mathbf{x})$ .
- Example:  $\psi(\mathbf{x}) = [x_1, x_2, x_1x_2]^T$ .

# Overcoming Limits: Feature Maps



- We can solve **XOR** by mapping input to a higher-dimensional space  $\psi(\mathbf{x})$ .
- Example:  $\psi(\mathbf{x}) = [x_1, x_2, x_1x_2]^T$ .

| Input |       | Feature Space |       |          | Tgt |
|-------|-------|---------------|-------|----------|-----|
| $x_1$ | $x_2$ | $x_1$         | $x_2$ | $\psi_3$ | $t$ |
| 0     | 0     | 0             | 0     | 0        | 0   |
| 0     | 1     | 0             | 1     | 0        | 1   |
| 1     | 0     | 1             | 0     | 0        | 1   |
| 1     | 1     | 1             | 1     | 1        | 0   |

**Insight:** In 3D, the **Red** points are *above* the separator, while **Blue** points are *below* it.

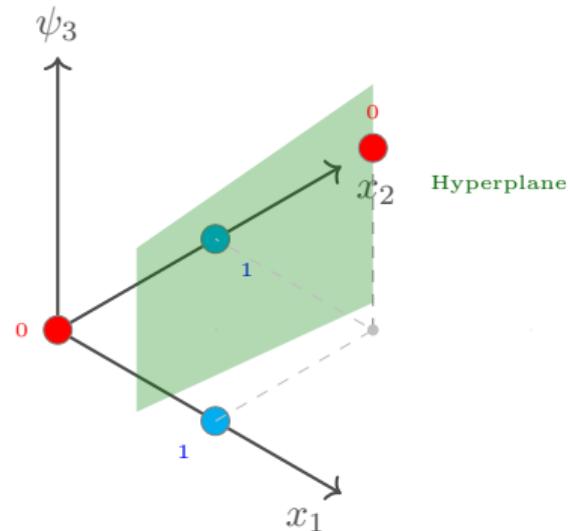
# Overcoming Limits: Feature Maps



- We can solve **XOR** by mapping input to a higher-dimensional space  $\psi(\mathbf{x})$ .
- Example:  $\psi(\mathbf{x}) = [x_1, x_2, x_1x_2]^T$ .

| Input |       | Feature Space |       |          | Tgt |
|-------|-------|---------------|-------|----------|-----|
| $x_1$ | $x_2$ | $x_1$         | $x_2$ | $\psi_3$ | $t$ |
| 0     | 0     | 0             | 0     | 0        | 0   |
| 0     | 1     | 0             | 1     | 0        | 1   |
| 1     | 0     | 1             | 0     | 0        | 1   |
| 1     | 1     | 1             | 1     | 1        | 0   |

**Insight:** In 3D, the Red points are *above* the separator, while Blue points are *below* it.

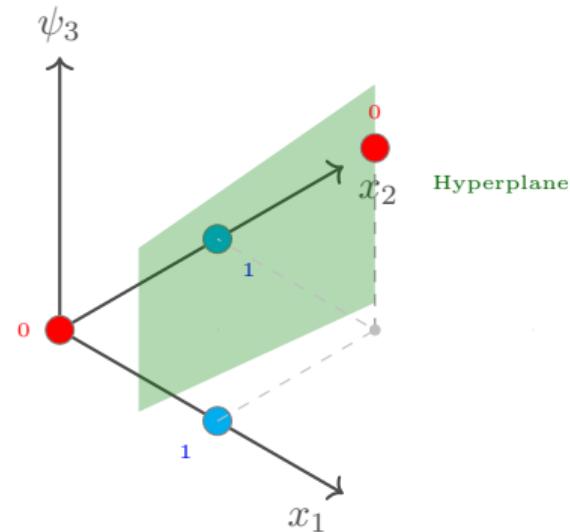


# Overcoming Limits: Feature Maps



- We can solve **XOR** by mapping input to a higher-dimensional space  $\psi(\mathbf{x})$ .
- Example:  $\psi(\mathbf{x}) = [x_1, x_2, x_1x_2]^T$ .

| Input |       | Feature Space |       |          | Tgt |
|-------|-------|---------------|-------|----------|-----|
| $x_1$ | $x_2$ | $x_1$         | $x_2$ | $\psi_3$ | $t$ |
| 0     | 0     | 0             | 0     | 0        | 0   |
| 0     | 1     | 0             | 1     | 0        | 1   |
| 1     | 0     | 1             | 0     | 0        | 1   |
| 1     | 1     | 1             | 1     | 1        | 0   |



**Insight:** In 3D, the **Red** points are *above* the separator, while **Blue** points are *below* it.

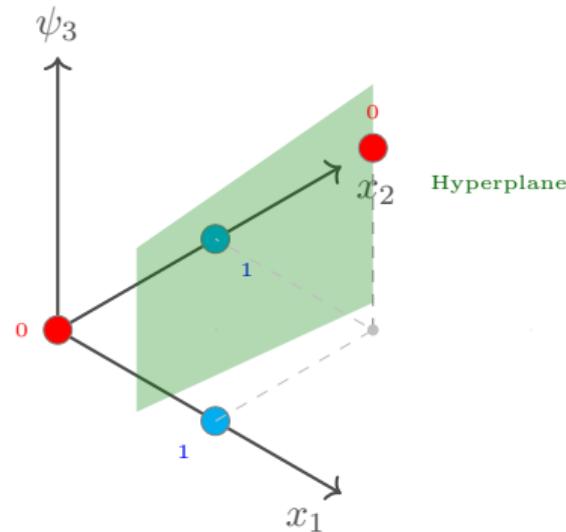
- In this new space, the problem is linearly separable.

# Overcoming Limits: Feature Maps



- We can solve **XOR** by mapping input to a higher-dimensional space  $\psi(\mathbf{x})$ .
- Example:  $\psi(\mathbf{x}) = [x_1, x_2, x_1x_2]^T$ .

| Input |       | Feature Space |       |          | Tgt |
|-------|-------|---------------|-------|----------|-----|
| $x_1$ | $x_2$ | $x_1$         | $x_2$ | $\psi_3$ | $t$ |
| 0     | 0     | 0             | 0     | 0        | 0   |
| 0     | 1     | 0             | 1     | 0        | 1   |
| 1     | 0     | 1             | 0     | 0        | 1   |
| 1     | 1     | 1             | 1     | 1        | 0   |



**Insight:** In 3D, the **Red** points are *above* the separator, while **Blue** points are *below* it.

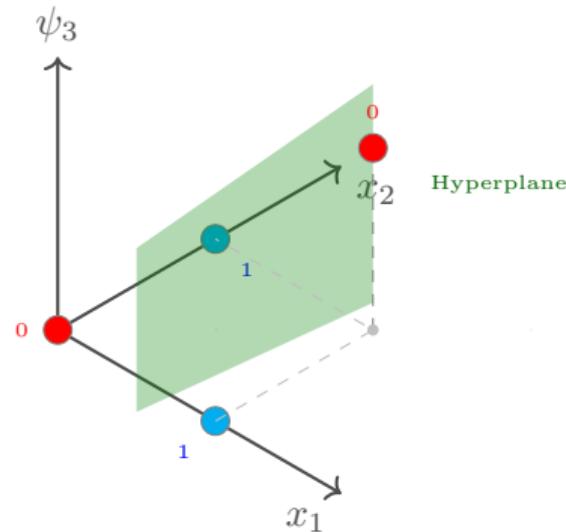
- In this new space, the problem is linearly separable.
- **Problem:** Hand-engineering features  $\psi(\mathbf{x})$  is difficult and domain-specific (e.g., for pixels in images).

# Overcoming Limits: Feature Maps



- We can solve **XOR** by mapping input to a higher-dimensional space  $\psi(\mathbf{x})$ .
- Example:  $\psi(\mathbf{x}) = [x_1, x_2, x_1x_2]^T$ .

| Input |       | Feature Space |       |          | Tgt |
|-------|-------|---------------|-------|----------|-----|
| $x_1$ | $x_2$ | $x_1$         | $x_2$ | $\psi_3$ | $t$ |
| 0     | 0     | 0             | 0     | 0        | 0   |
| 0     | 1     | 0             | 1     | 0        | 1   |
| 1     | 0     | 1             | 0     | 0        | 1   |
| 1     | 1     | 1             | 1     | 1        | 0   |



**Insight:** In 3D, the **Red** points are *above* the separator, while **Blue** points are *below* it.

- In this new space, the problem is linearly separable.
- **Problem:** Hand-engineering features  $\psi(\mathbf{x})$  is difficult and domain-specific (e.g., for pixels in images).
- **Solution:** Learn the non-linear mapping  $\psi$  automatically! → **Multiple Layer of Neural Networks**.

# Contents



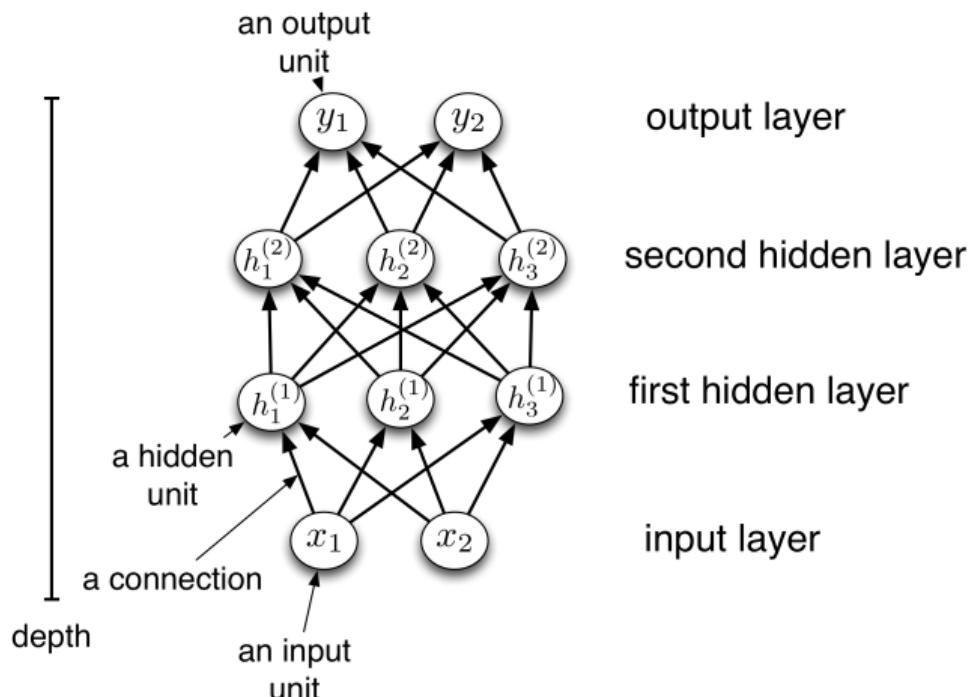
- ▶ Limitations of Linear Models
- ▶ Multilayer Perceptrons (MLPs)

- ▶ MLPs in the Modern Era

# Multilayer Perceptrons (MLPs)



- An MLP is a **Directed Acyclic Graph (DAG)** of units, organized in layers.
- Also known as a **Feed-Forward Network**.
- **Architecture:**
  1. **Input Layer:** Raw data  $x$
  2. **Hidden Layers:** Learn intermediate representations  $h^{(l)}$
  3. **Output Layer:** Predictions  $y$



# Fully Connected Layers



- In the simplest case, all input units are connected to all output units. We call this a **fully connected layer**.

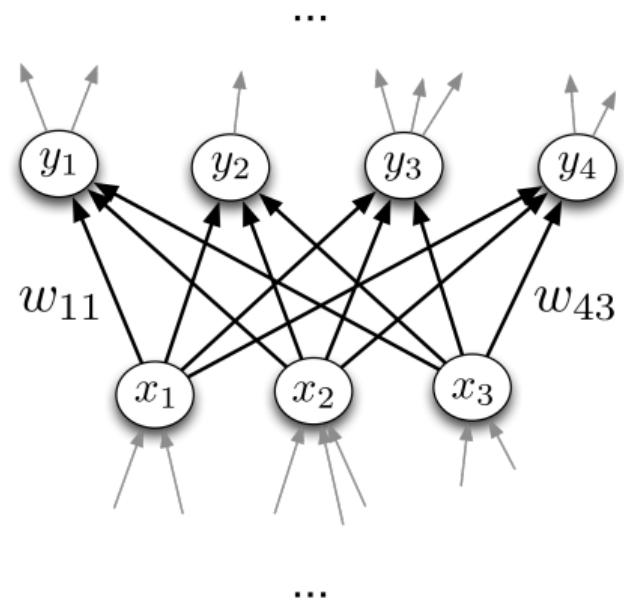
## Layer $l$ Computation

For every layer  $l = 1 \dots L$ :

$$\mathbf{z}^{(l)} = \mathbf{W}^{(l)} \mathbf{h}^{(l-1)} + \mathbf{b}^{(l)} \quad (\text{Linear})$$

$$\mathbf{h}^{(l)} = \phi(\mathbf{z}^{(l)}) \quad (\text{Non-linear})$$

- $\mathbf{W}^{(l)}$ : Weight matrix (learnable parameters).
- $\mathbf{b}^{(l)}$ : Bias vector.
- $\phi$ : Activation function (applied element-wise).



# Why Non-linearities?



- What if we stack linear layers without  $\phi$ ?

$$\mathbf{y} = \mathbf{W}_2(\mathbf{W}_1\mathbf{x}) = (\mathbf{W}_2\mathbf{W}_1)\mathbf{x} = \mathbf{W}'\mathbf{x}$$

# Why Non-linearities?



- What if we stack linear layers without  $\phi$ ?

$$\mathbf{y} = \mathbf{W}_2(\mathbf{W}_1\mathbf{x}) = (\mathbf{W}_2\mathbf{W}_1)\mathbf{x} = \mathbf{W}'\mathbf{x}$$

## Crucial Insight

- The composition of linear functions is just another linear function.

# Why Non-linearities?



- What if we stack linear layers without  $\phi$ ?

$$\mathbf{y} = \mathbf{W}_2(\mathbf{W}_1\mathbf{x}) = (\mathbf{W}_2\mathbf{W}_1)\mathbf{x} = \mathbf{W}'\mathbf{x}$$

## Crucial Insight

- The composition of linear functions is just another linear function.
- Deep linear networks have no more expressive power than a single linear layer!

# Why Non-linearities?



- What if we stack linear layers without  $\phi$ ?

$$\mathbf{y} = \mathbf{W}_2(\mathbf{W}_1\mathbf{x}) = (\mathbf{W}_2\mathbf{W}_1)\mathbf{x} = \mathbf{W}'\mathbf{x}$$

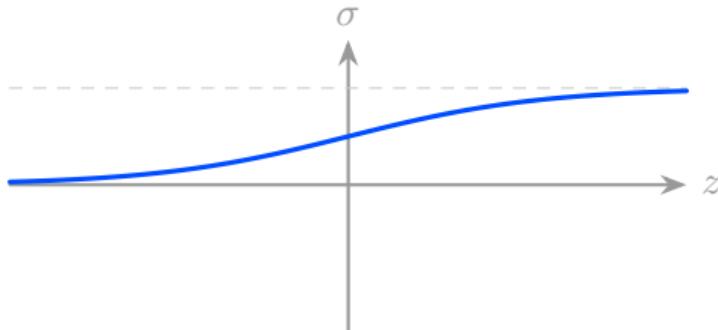
## Crucial Insight

- The composition of linear functions is just another linear function.
- Deep linear networks have no more expressive power than a single linear layer!
- **Activation functions** break linearity, allowing us to approximate curved boundaries and complex manifolds.

# Common Activation Functions (1)



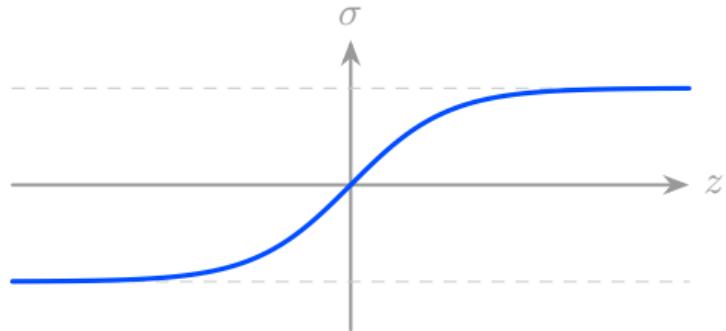
Sigmoid



$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

- + Interpretation: Probability
- Vanishing gradients
- Output not zero-centered

Tanh



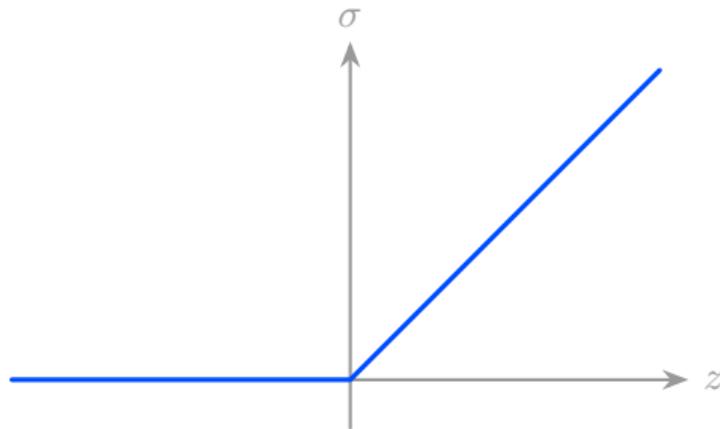
$$\sigma(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$$

- + Zero-centered output
- + Stronger gradients than Sigmoid
- Still saturates (vanishing grad)

# Common Activation Functions (2)



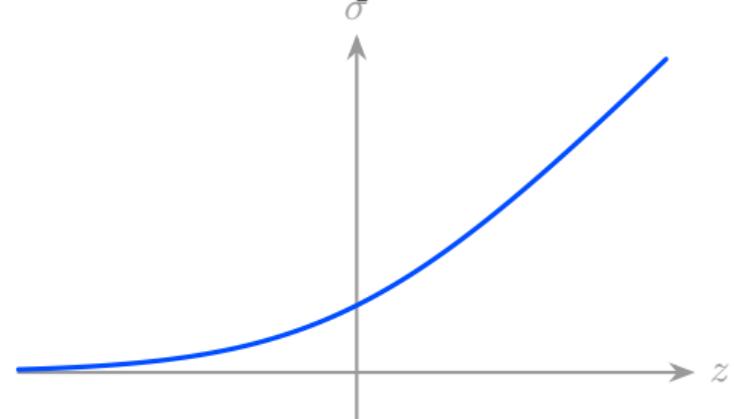
ReLU



$$\sigma(z) = \max(0, z)$$

- + Efficient; Sparse activations
- + No vanishing grad ( $z > 0$ )
- “Dead” neurons ( $z < 0$ )

Softplus



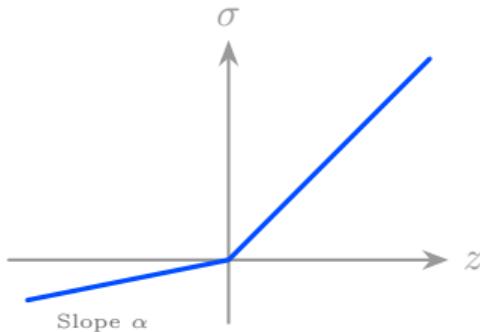
$$\sigma(z) = \ln(1 + e^z)$$

- + Smooth approx. of ReLU
- + Differentiable everywhere
- Computationally slower

# Modern Activation Functions (3)

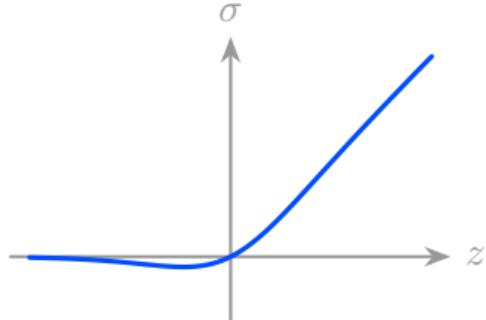


Leaky ReLU



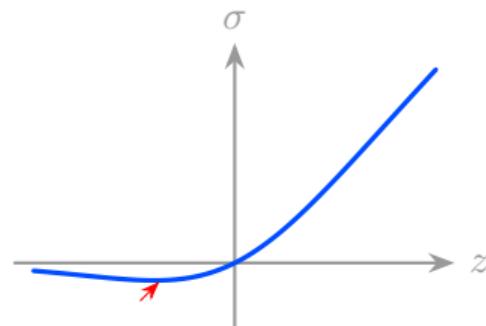
$$\sigma(z) = \max(\alpha z, z)$$

GELU



$$\sigma(z) \approx z\Phi(z)$$

Swish / SiLU



+ Fixes “Dead ReLU”  
Note:  $\alpha$  is small  
Use: GANs, Deep MLPs

Note:  $\Phi(z)$  is CDF of Normal  
Core: GPT, BERT, ViT  
+ Probabilistic  
+ Optimization aid

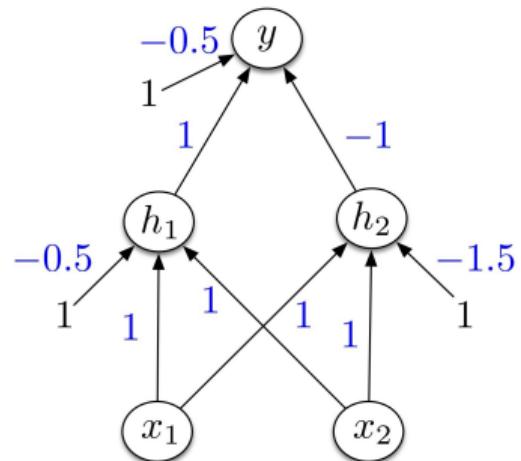
Note:  $\sigma_s(z)$  is Sigmoid  
Core: EfficientNet, YOLO  
+ Non-monotonic dip  
+ Self-gating

# Solving XOR with an MLP



We can design a network that computes XOR *exactly* using hard threshold function:

$$\sigma(z) = \begin{cases} 0 & \text{if } z < 0 \\ 1 & \text{if } z \geq 0 \end{cases}$$



## Logic Circuits

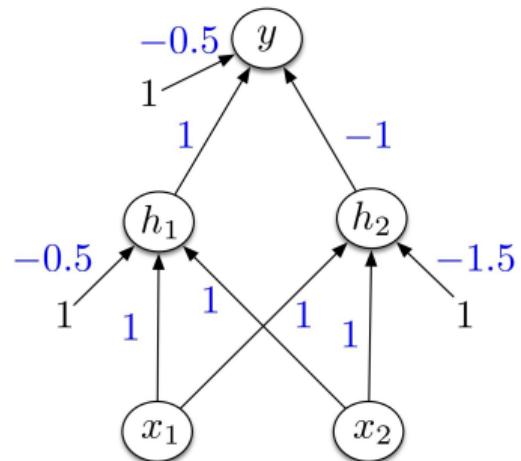
- $h_1$  acts as **AND**= $x_1 \wedge x_2$ :  $\sigma(x_1 + x_2 - 1.5)$

# Solving XOR with an MLP



We can design a network that computes XOR *exactly* using hard threshold function:

$$\sigma(z) = \begin{cases} 0 & \text{if } z < 0 \\ 1 & \text{if } z \geq 0 \end{cases}$$



## Logic Circuits

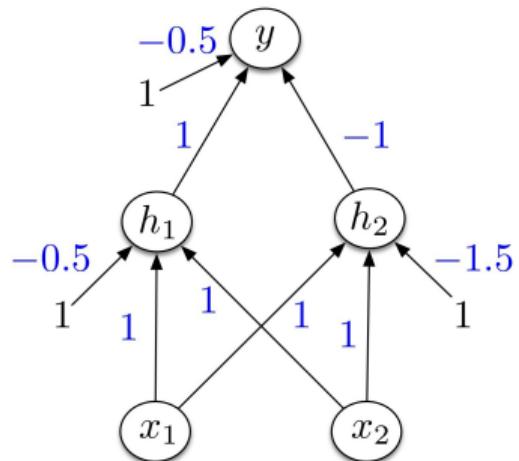
- $h_1$  acts as **AND**= $x_1 \wedge x_2$ :  $\sigma(x_1 + x_2 - 1.5)$
- $h_2$  acts as **OR**= $x_1 \vee x_2$ :  $\sigma(x_1 + x_2 - 0.5)$

# Solving XOR with an MLP



We can design a network that computes XOR *exactly* using hard threshold function:

$$\sigma(z) = \begin{cases} 0 & \text{if } z < 0 \\ 1 & \text{if } z \geq 0 \end{cases}$$



## Logic Circuits

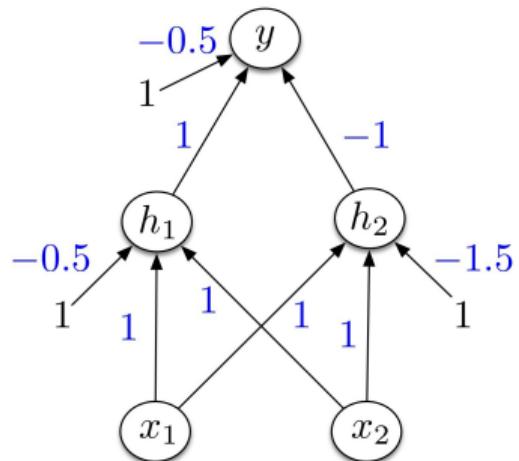
- $h_1$  acts as **AND**= $x_1 \wedge x_2$ :  $\sigma(x_1 + x_2 - 1.5)$
- $h_2$  acts as **OR**= $x_1 \vee x_2$ :  $\sigma(x_1 + x_2 - 0.5)$
- $y$  acts as **XOR**= $(x_1 \vee x_2) \wedge \neg(x_1 \wedge x_2)$ :  
 $\sigma(h_2 - h_1 - 0.5)$

# Solving XOR with an MLP



We can design a network that computes XOR *exactly* using hard threshold function:

$$\sigma(z) = \begin{cases} 0 & \text{if } z < 0 \\ 1 & \text{if } z \geq 0 \end{cases}$$



## Logic Circuits

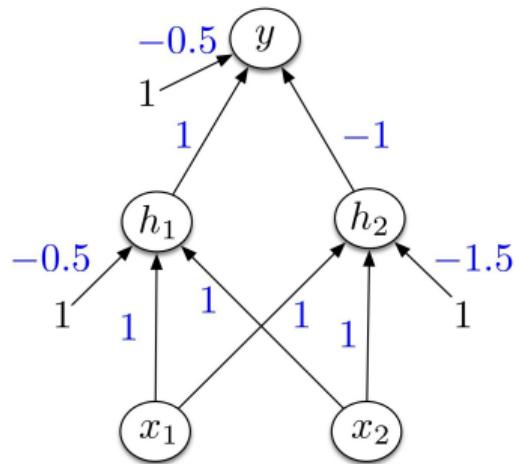
- $h_1$  acts as **AND**= $x_1 \wedge x_2$ :  $\sigma(x_1 + x_2 - 1.5)$
- $h_2$  acts as **OR**= $x_1 \vee x_2$ :  $\sigma(x_1 + x_2 - 0.5)$
- $y$  acts as **XOR**= $(x_1 \vee x_2) \wedge \neg(x_1 \wedge x_2)$ :  
 $\sigma(h_2 - h_1 - 0.5)$

# Solving XOR with an MLP



We can design a network that computes XOR *exactly* using hard threshold function:

$$\sigma(z) = \begin{cases} 0 & \text{if } z < 0 \\ 1 & \text{if } z \geq 0 \end{cases}$$



## Logic Circuits

- $h_1$  acts as **AND**= $x_1 \wedge x_2$ :  $\sigma(x_1 + x_2 - 1.5)$
- $h_2$  acts as **OR**= $x_1 \vee x_2$ :  $\sigma(x_1 + x_2 - 0.5)$
- $y$  acts as **XOR**= $(x_1 \vee x_2) \wedge \neg(x_1 \wedge x_2)$ :  
 $\sigma(h_2 - h_1 - 0.5)$

Trace for input (1, 1):

$$h_1 = \sigma(1 + 1 - 1.5) = \sigma(0.5) = 1$$

$$h_2 = \sigma(1 + 1 - 0.5) = \sigma(1.5) = 1$$

$$y = \sigma(1 - 1 - 0.5) = \sigma(-0.5) = 0$$

# MLP as Feature Learning



Each layer computes a function, so the network computes a composition of functions:

$$\mathbf{h}^{(1)} = f^{(1)}(\mathbf{x})$$

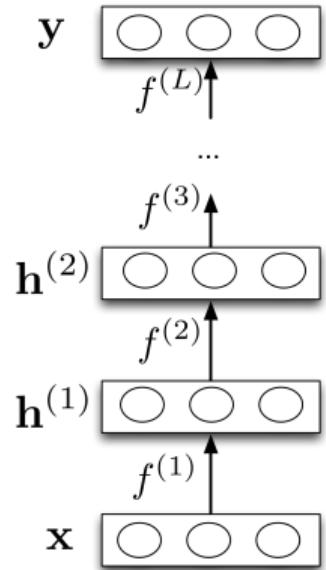
$$\mathbf{h}^{(2)} = f^{(2)}(\mathbf{h}^{(1)})$$

⋮

$$\mathbf{y} = f^{(L)}(\mathbf{h}^{(L-1)})$$

Or more generally:  $\mathbf{y} = f^{(L)} \circ f^{(L-1)} \circ \dots \circ f^{(1)}(\mathbf{x})$

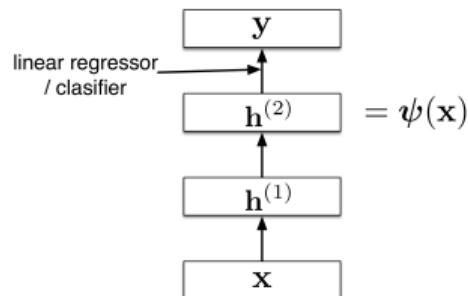
- Each layer transforms the input into a new, more abstract representation.
- The final layer produces the output based on these learned features.



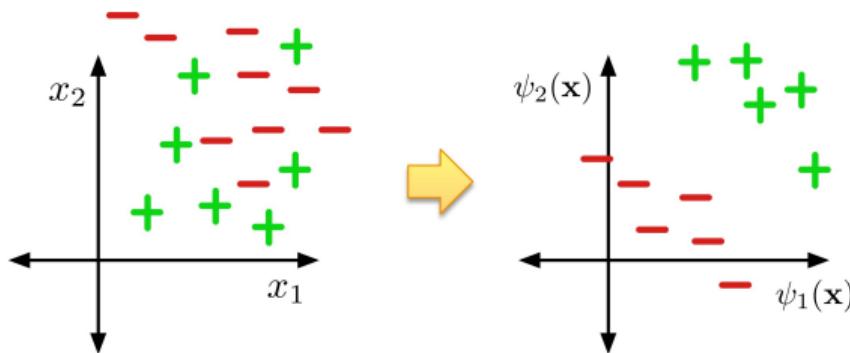
# MLP as Feature Learning



- Neural nets can be viewed as a way of learning features:



- The goal:

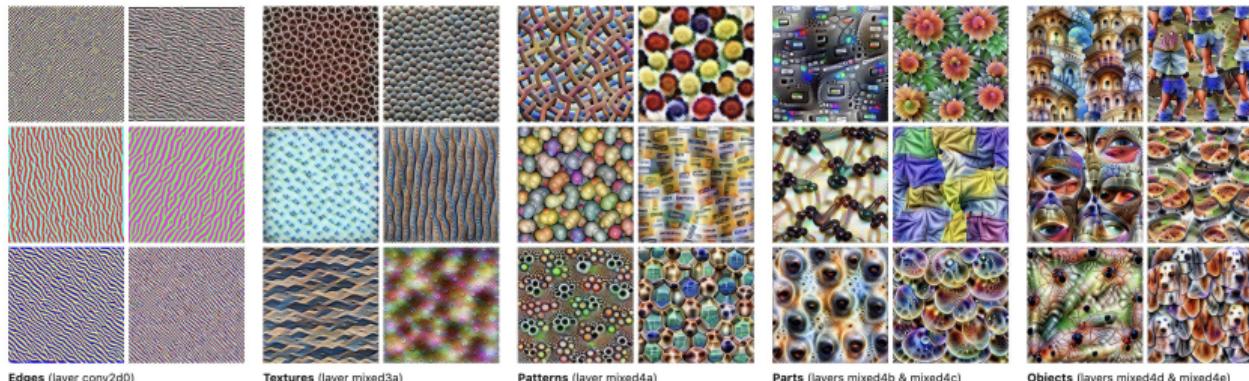


# MLP as Feature Learning: Image Classification



**Idea:** When multiple layers are stacked, the network learns hierarchical features:

- **Lower layers:** Detect simple patterns such as edges.
- **Middle layers:** Combine edges into parts.
- **Higher layers:** Assemble parts into objects.



Example: GoogLeNet trained on ImageNet reveals how features evolve from edges to complex objects using *DeepDream*.

# Designing the Output Layer



The choice of output function  $\phi_{out}$  and Loss function  $\mathcal{L}$  depends on the task:

| Task          | Target $t$            | Output Unit         | Loss Function        |
|---------------|-----------------------|---------------------|----------------------|
| Regression    | $t \in \mathbb{R}$    | Linear (Identity)   | Squared Error        |
| Binary Class. | $t \in \{0, 1\}$      | Sigmoid $\sigma(z)$ | Binary Cross-Entropy |
| Multi-class   | $t \in \{1 \dots K\}$ | <b>Softmax</b>      | <b>Cross-Entropy</b> |

## Universal Approximation Theorem

A feed-forward neural network with a single hidden layer containing a finite number of neurons can approximate any continuous function on compact subsets of  $\mathbb{R}^n$  to arbitrary accuracy.

## Universal Approximation Theorem

A feed-forward neural network with a single hidden layer containing a finite number of neurons can approximate any continuous function on compact subsets of  $\mathbb{R}^n$  to arbitrary accuracy.

- **Implication:** Neural networks are universal function approximators.
- **Caveat:** The theorem says a network *exists*, but doesn't say how to find it (optimization) or how large it needs to be (efficiency).
- Various activation sufficient for approximation.
- We often prefer **deep** (more layers) networks over wide ones for better generalization and parameter efficiency.

# Contents



- ▶ Limitations of Linear Models
- ▶ Multilayer Perceptrons (MLPs)
- ▶ MLPs in the Modern Era

# Myth vs. Reality



## Myth

“Deep Learning replaced MLPs with ConvNets and Transformers.”

# Myth vs. Reality



## Myth

“Deep Learning replaced MLPs with ConvNets and Transformers.”

- **Reality: MLPs never left.** They just changed their name.
- In modern architectures (Transformers, NeRFs), MLPs typically constitute **2/3 of the total parameters.**
- In NeRF, the entire scene is an MLP.

Self-Attention

MLP / FFN  
(Feed Forward)

Layer Norm

← In LLMs (GPT, Llama), this block runs on *every token* and consumes the majority of compute!

# Role 1: The “Channel Mixer” (Transformers)



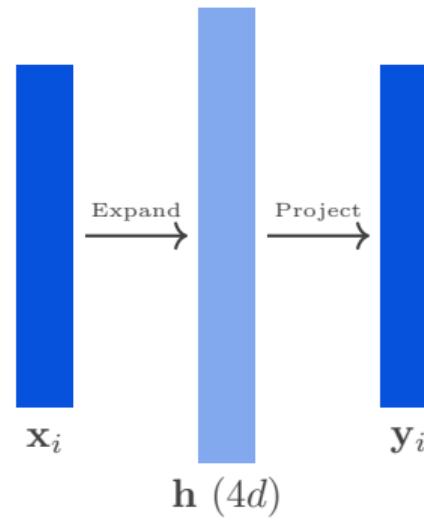
In Transformers (e.g., ChatGPT), we distinguish two types of operations:

## 1. Token Mixing (Attention)

- Mixes information *between* different words in the sequence.
- “Looking at other words”

## 2. Channel Mixing (MLP)

- Mixes information *within* a single word embedding.
- Processing the features extracted by attention.
- Applied **position-wise** (independently to every token).



$$\text{FFN}(x) = \phi(x\mathbf{W}_1 + \mathbf{b}_1)\mathbf{W}_2 + \mathbf{b}_2$$



## The Hypothesis (Geva et al., 2021)

Transformer MLPs function as **associative memories**, storing factual knowledge in their weights.

### The Mechanism:

#### 1. Pattern Detection (The Key)

- The input  $\mathbf{x}$  acts as a **query**.
- It scans the rows of  $\mathbf{W}_1$  to find a matching pattern (e.g., “Capital” + “France”).
- *Math:*  $k = \sigma(\mathbf{x} \cdot \mathbf{W}_1^T)$

# Role 2: Knowledge Storage (Key-Value Memories)



## The Hypothesis (Geva et al., 2021)

Transformer MLPs function as **associative memories**, storing factual knowledge in their weights.

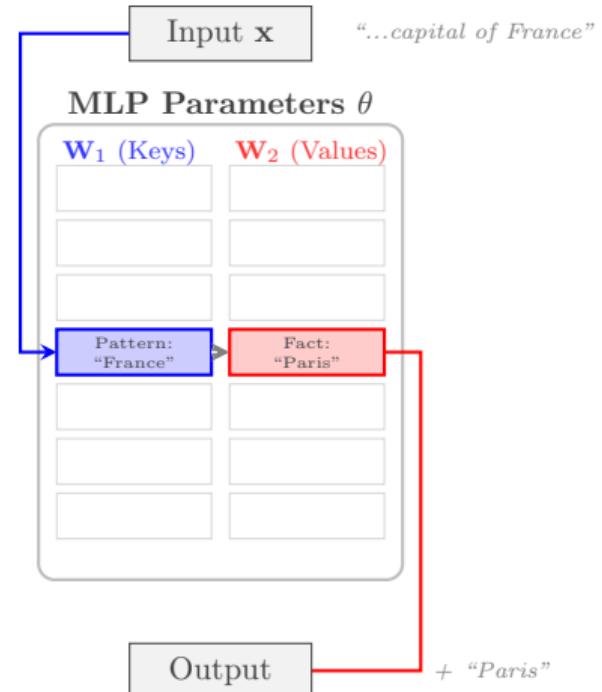
## The Mechanism:

### 1. Pattern Detection (The Key)

- The input  $\mathbf{x}$  acts as a **query**.
- It scans the rows of  $\mathbf{W}_1$  to find a matching pattern (e.g., “Capital” + “France”).
- *Math:*  $k = \sigma(\mathbf{x} \cdot \mathbf{W}_1^T)$

### 2. Fact Retrieval (The Value)

- The active neuron “unlocks” the corresponding row in  $\mathbf{W}_2$ .
- The network retrieves the stored vector (e.g., “Paris”).
- *Math:*  $\mathbf{y} = k \cdot \mathbf{W}_2$



# Role 3: Implicit Neural Representations (INRs)

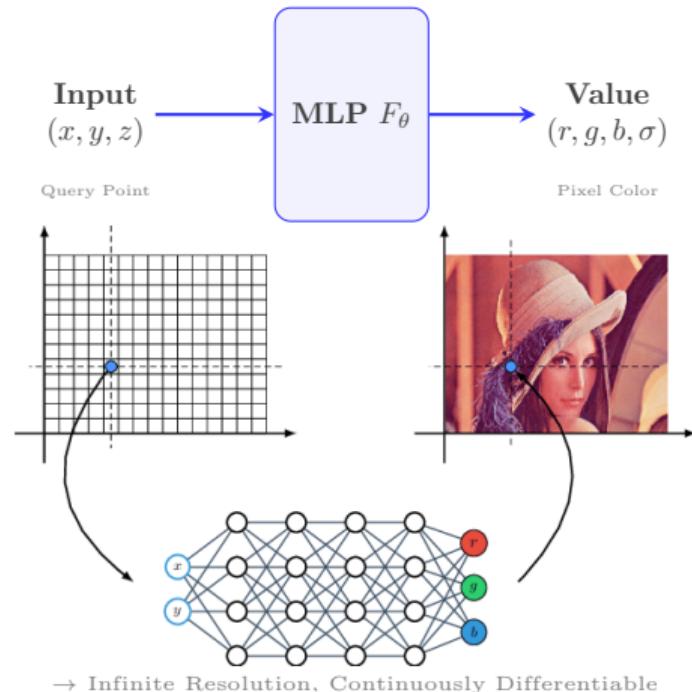


## Concept: Functions as Data

MLPs are not just for *processing* data; they can be the data.

### Coordinate-based MLPs:

- Instead of storing data in a discrete grid (pixels, voxels), we represent it as a continuous function  $F_\theta$ .



# Role 3: Implicit Neural Representations (INRs)



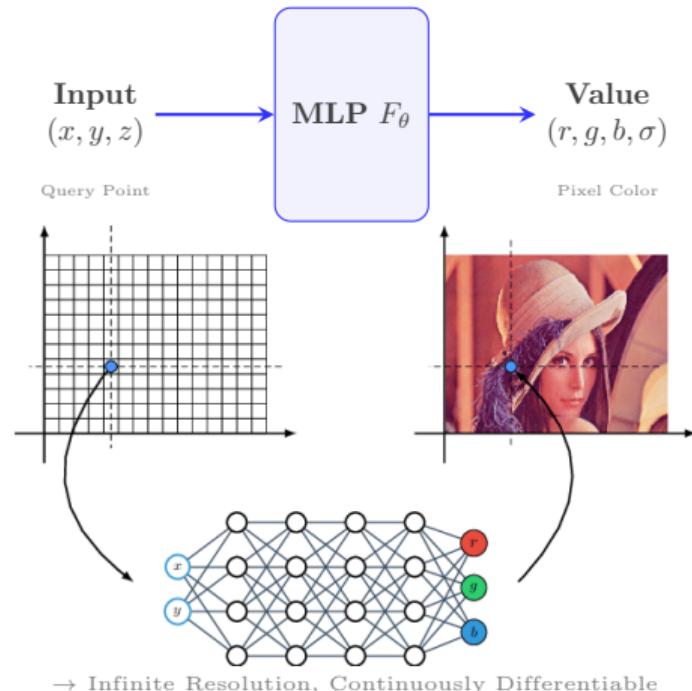
## Concept: Functions as Data

MLPs are not just for *processing* data; they can be the data.

### Coordinate-based MLPs:

- Instead of storing data in a discrete grid (pixels, voxels), we represent it as a continuous function  $F_\theta$ .
- $F_\theta$  as a **coordinate-wise look-up table**, mapping a position  $\mathbf{x}$  to a value:

$$F_\theta(\mathbf{x}) \approx \text{Signal}(\mathbf{x})$$



# Role 3: Implicit Neural Representations (INRs)



## Concept: Functions as Data

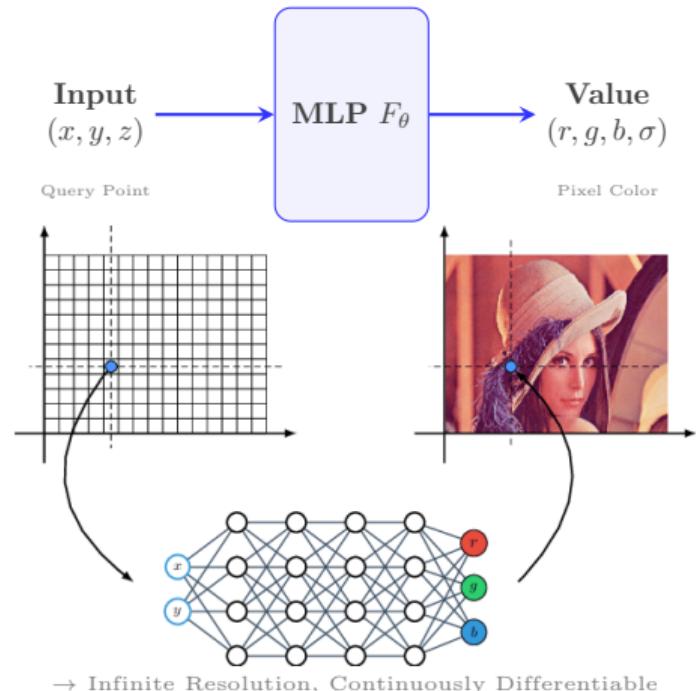
MLPs are not just for *processing* data; they can be the data.

### Coordinate-based MLPs:

- Instead of storing data in a discrete grid (pixels, voxels), we represent it as a continuous function  $F_\theta$ .
- $F_\theta$  as a **coordinate-wise look-up table**, mapping a position  $\mathbf{x}$  to a value:

$$F_\theta(\mathbf{x}) \approx \text{Signal}(\mathbf{x})$$

- The network weights  $\theta$  act as the **compression storage**.



# Role 3: Implicit Neural Representations (INRs)



## Concept: Functions as Data

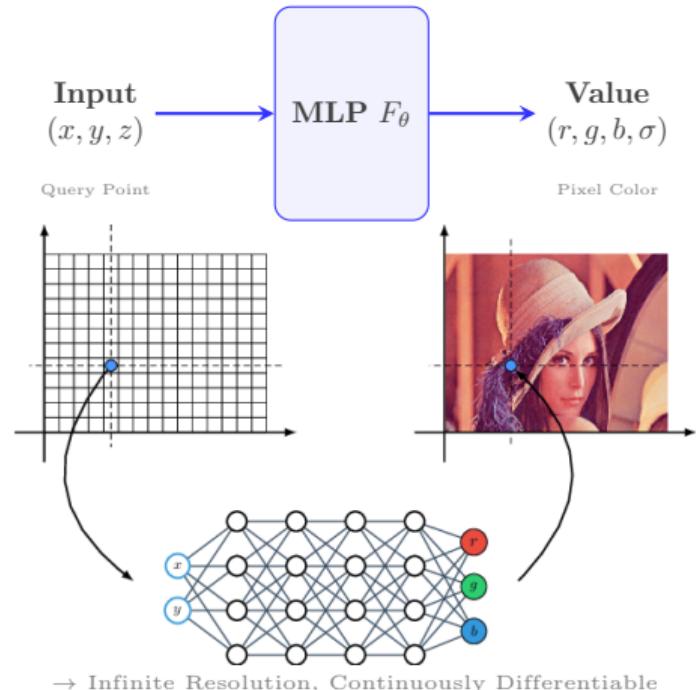
MLPs are not just for *processing* data; they can be the data.

### Coordinate-based MLPs:

- Instead of storing data in a discrete grid (pixels, voxels), we represent it as a continuous function  $F_\theta$ .
- $F_\theta$  as a **coordinate-wise look-up table**, mapping a position  $\mathbf{x}$  to a value:

$$F_\theta(\mathbf{x}) \approx \text{Signal}(\mathbf{x})$$

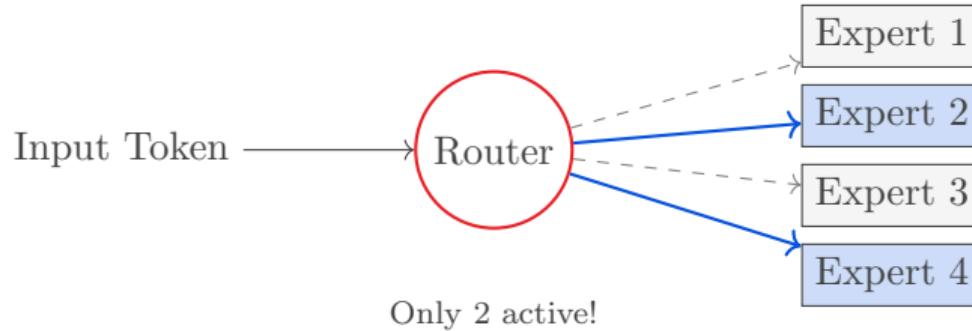
- The network weights  $\theta$  act as the **compression storage**.
- **Example:** neural radiance fields (NeRFs) for 3D rendering.



## Role 4: Mixture of Experts (MoE)



- As models grow, dense MLPs become too expensive to compute.
- **Solution: Mixture of Experts (MoE).**
- Replace one giant MLP with  $N$  smaller “expert” MLPs.
- A **Router** network selects only the top- $k$  (e.g., 2) experts for each token.



**Impact:** Massive capacity (params) with low inference cost (FLOPs). Used in **Mixtral**, **GPT-4**.

# Summary



1. **Linear Models** are limited (XOR problem, no compositionality).
2. **MLPs** introduce non-linearity via activation functions ( $\sigma$ ), allowing universal approximation.
3. **Deep Learning** leverages depth to learn hierarchical features efficiently.
4. **Modern Usage:** MLPs are the workhorse of Transformers (Channel Mixing) and NeRFs.

# Next Up: Backpropagation



## The Big Question

The Universal Approximation Theorem guarantees that there exists a set of weights  $\mathbf{W}^*$  that approximates our function  $f(x)$ .

**But how do we find  $\mathbf{W}^*$ ?**

# Next Up: Backpropagation



## The Big Question

The Universal Approximation Theorem guarantees that there exists a set of weights  $\mathbf{W}^*$  that approximates our function  $f(x)$ .

But how do we find  $\mathbf{W}^*$ ?

We treat this as an **Optimization Problem**:

1. Define a **Loss Function**  $\mathcal{L}(\mathbf{y}, \mathbf{t})$  that measures “badness”.
2. Minimize this loss w.r.t parameters  $\theta = \{\mathbf{W}, \mathbf{b}\}$ .
3. Use **Gradient Descent**:  $\theta \leftarrow \theta - \eta \nabla_{\theta} \mathcal{L}$ .

# Next Up: Backpropagation



## The Big Question

The Universal Approximation Theorem guarantees that there exists a set of weights  $\mathbf{W}^*$  that approximates our function  $f(x)$ .

But how do we find  $\mathbf{W}^*$ ?

We treat this as an **Optimization Problem**:

1. Define a **Loss Function**  $\mathcal{L}(\mathbf{y}, \mathbf{t})$  that measures “badness”.
2. Minimize this loss w.r.t parameters  $\theta = \{\mathbf{W}, \mathbf{b}\}$ .
3. Use **Gradient Descent**:  $\theta \leftarrow \theta - \eta \nabla_{\theta} \mathcal{L}$ .

To do this, we will invent **Backpropagation** for neural nets.

See how to do backprop in next part!